

House Price Prediction System

Yuen Yee Lo, PhD

Data Science Intensive Capstone Project, May 2022



Contents

- Introduction
- Data
- EDA
- Modelling
- Results and Analysis
- Conclusion



Why important?

- Help people buy a house
- Know the price range in the future
- Plan their finance
- Beneficial for property investors
- Know the trend of housing prices in a certain location

Who May Care?

House buyer



Real estates




Data


- 1461 entries
- 81 explanatory variables describing (almost) every aspect
 - 36 numerical data
 - 43 category data
- Residential homes in Ames, Iowa
- 2006-2010

EDA

Sale Price distribution


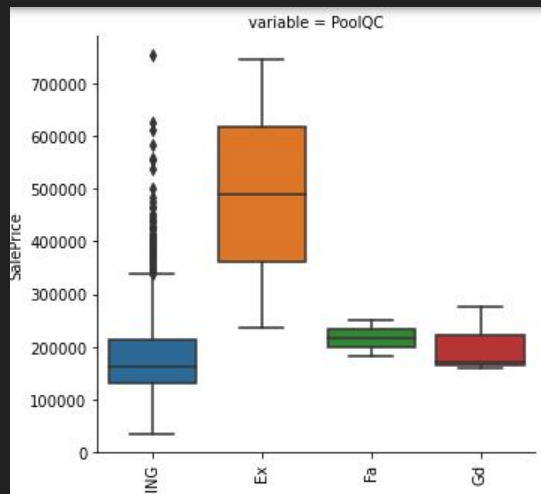



EDA: Category data : Disparity



EDA:

- Some categories seem to more diverse
- Neighborhood has big impact on house prices
- Having pool on property seems to improve price
- Partial SaleCondition has impact



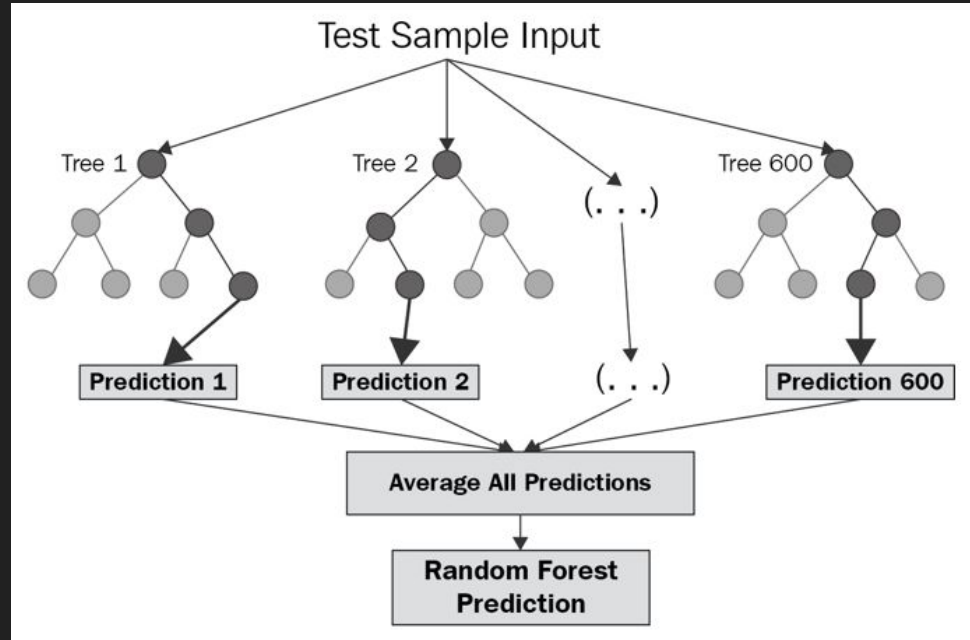
Modelling

Modelling: Baseline Linear Regression

- Linear approach for modelling the relationship between a scalar response and one or more explanatory variables
- Simple model
- Only numerical data

Modelling: Random Forest Regression

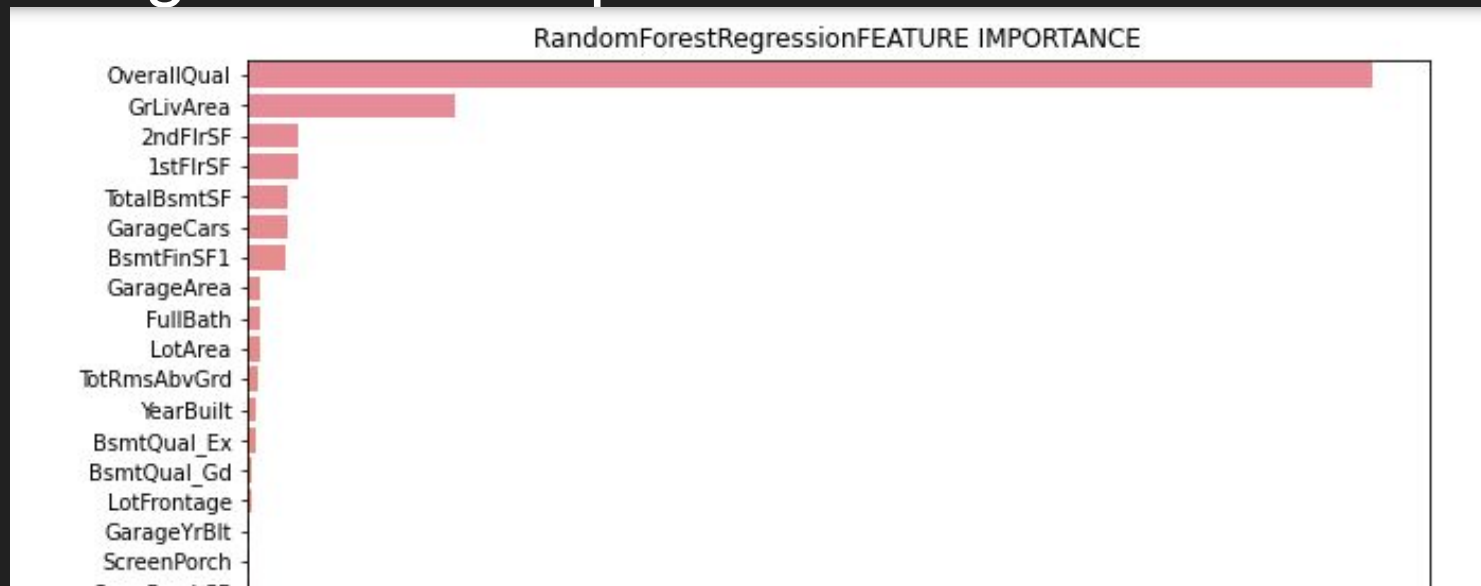
- Ensemble technique
- Regression and classification
- Bootstrap multiple DTs
- Reduce overfitting



Modelling: Feature Importance


- Random forest can give us insight about the feature importance
- Reduce number of features
- Reduces the complexity of a model
- Easier to interpret
- Improves the accuracy if the right subset is chosen

Modelling: Feature Importance



- Feature “OverallQual” is the most important features
- Top10 important features are numerical data
- Category data are not significant

OverallQual vs. SalePrice




Modelling: **Gradient Boosting Regression**

- Trains many models in a gradual, additive and sequential manner
- The major difference between AdaBoost and Gradient Boosting Algorithm is how the two algorithms identify the shortcomings of weak learners (eg. decision trees)
- Loss function would be based off the error between true and predicted house prices

Gradient Boosting

Gradient Boosting Model

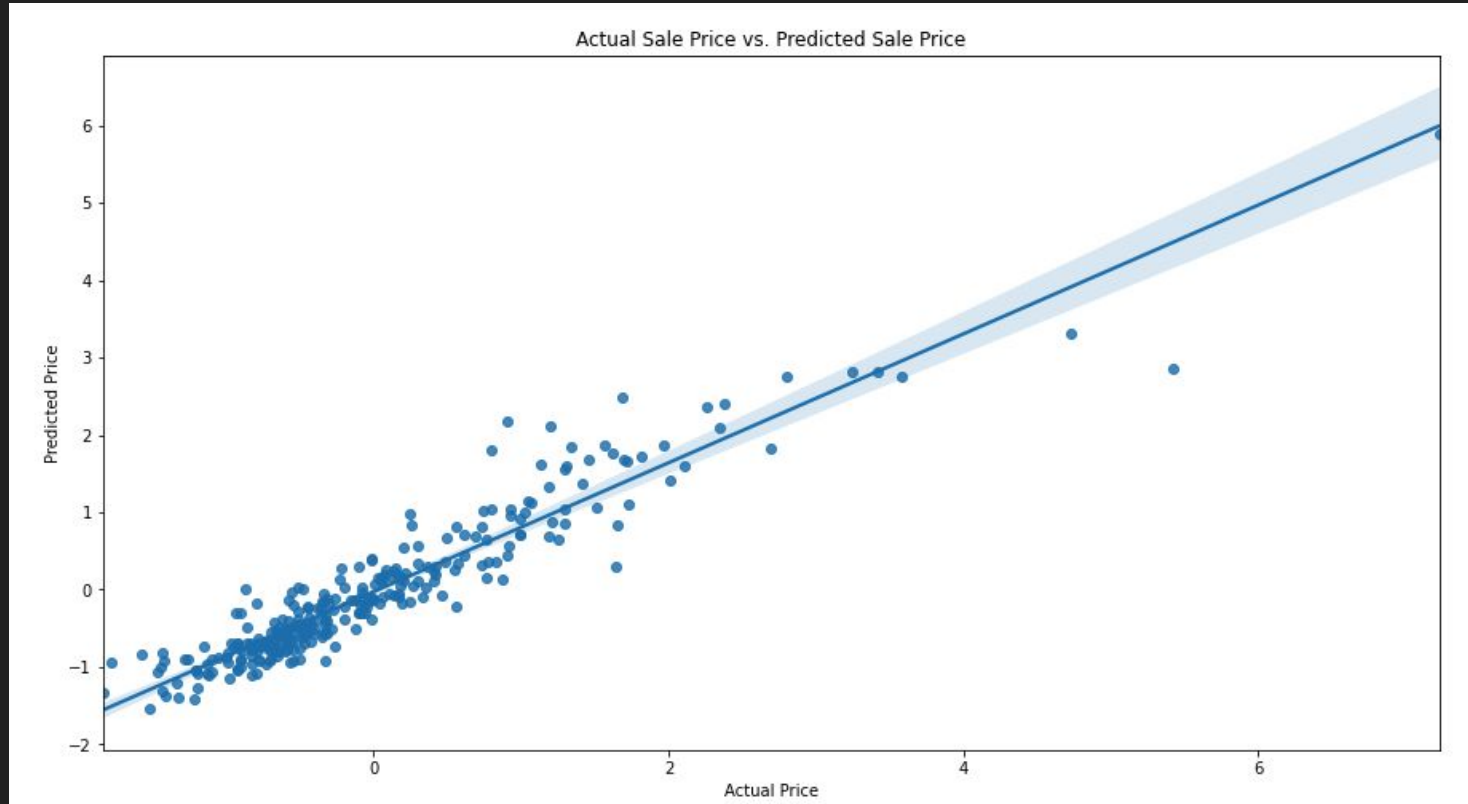


Results:

- MAE, MSE, RMSE for result matrix
- Gradient boosting performs the best
- Reduce features can improve the performance

	MAE	MSE	RMSE
Linear Regression	0.289	0.215	0.463
Random Forest	0.288	0.199	0.447
Random Forest Top10 features	0.288	0.187	0.433
Gradient Boosting	0.241	0.133	0.365

Predicted housing price (test data)



Conclusions

- Built a baseline model using linear regression
- Compare to the random forest regression and gradient boosting regression
- Feature selection from the features importance of RF
- Categories data show no significant impact
- Gradient boosting regression performs the best
- Further improvement:
 - parameter tuning
 - deep learning (if we have more data)

Thank You

Question?

Contact

Yuen Yee Lo

Email: yylo7775@gmail.com

<https://www.linkedin.com/in/yuen-yee-lo-4b74865/>

<https://github.com/yuenyeelo/>

<http://www.y2nlp.com>

Project directory: <https://github.com/yuenyeelo/springboard/tree/main/Capstone2>